

JONATHAN JIN

me@jonathanjin.in
linkedin.com/in/jinnovation
github.com/jinnovation
huggingface.co/jinnovation

SELECT WORK EXPERIENCE

Hinge

Staff Machine Learning Engineer

New York, NY, Apr 2024 – Current

Leading the democratization and platformization of AI at Hinge. Spearheading 0-to-1 development of agentic AI infrastructure. Previously defined, led, and delivered on traditional-ML model serving strategy from end-to-end.

- Tech lead for “User Traits,” an agentic system leveraging LLMs and VLMs to power enhanced user understanding. Defined the shared semantic layer as well as the foundational architecture (TensorZero, OpenTelemetry, Spark, Unity Catalog). Collaborated with development partners across Backend, Data, and AI disciplines to identify use cases and coordinate contributions.
- Laid the foundations of Hinge’s centralized agentic AI infrastructure. Successfully led and closed out build-or-buy evaluation of vendor/OSS solutions for evaluation and prompt engineering/management. Implemented platform based on OpenTelemetry, OpenInference, and Phoenix in collaboration with stakeholders across Backend, AI, and Data disciplines. Identified need for centralized LLM gateway solution to federate foundation model queries across model providers and query patterns. Led partnership with TensorZero team on open-source gateway and optimization solution.
- Defined and led the implementation and delivery of Hinge AI’s model serving strategy. Collaborated with applied AI/ML teams to coalesce existing approaches/solutions into cohesive platform based on MLflow, Ray, and Spark, reducing time-to-production for new models by ~50%. Defined strategic principles, focuses, and year-long roadmap in close collaboration with platform partner teams in Data and Engineering to achieve general availability in three quarters.

Humane Intelligence

Engineering Advisor

New York, NY, Feb 2024 – Aug 2024

“Fractional” tech lead and architect. Partner with the CEO to define and execute on engineering strategies. Contribute in hands-on capacity for technical proofs-of-concepts.

Spotify

Senior Machine Learning Engineer

New York, NY, Mar 2021 – Mar 2024

Senior member of the ML Platform org. Primarily focused on managed multi-tenant compute infrastructure for Kubeflow and Ray in GCP. Later shifted focus to AI/ML governance and artifact/experiment management solutions.

- Drove early-stage evaluation and implementation of centralized, multi-tenant Ray infrastructure based on KubeRay;
- Led development of Backstage-based model governance solution in preparation for the EU Digital Services Act;
- Defined general-availability roadmap for Kubeflow-based orchestration product, centered on SLOs and SRE best practices.

NVIDIA

Senior Systems Software Engineer, AI Infrastructure

New York, NY, Dec 2019 – Jan 2021

Worked on AI infrastructure for autonomous vehicle research and development.

Twitter

Machine Learning Software Engineer

New York, NY, Aug 2018 – Dec 2019

Member of Cortex, Twitter’s central ML platform organization. Worked on: workflow orchestration; experiment management.

EDUCATION

University of Chicago B.S. Computer Science, B.A Economics

2011 – 2015

SPEAKING

Data Council, “Trimming the Long Tail of Production Model Ownership at Hinge”

Apr 2025

Ray Summit, “Accelerating Time-to-Production for ML at Hinge”

Sep 2024

MLconf, “Empowering Traceable and Auditable ML in Production at Spotify with Hendrix”

Mar 2022

KubeCon + CloudNativeCon, “Scaling Kubeflow for Multi-tenancy at Spotify”

Oct 2021

SKILLS

- **AI, ML, and Data:** TensorZero, Phoenix, OpenInference, Ray, PyTorch, Spark, Flink, MLflow, Airflow, Kubeflow, TensorFlow
- **Infrastructure and Distributed Systems:** Kubernetes, Terraform, Helm, Prometheus, OpenTelemetry, gRPC, Docker

- **Enterprise Platforms:** Amazon Web Services (AWS), Google Cloud (GCP), Databricks, Arize
- **Programming Languages:** Python, Go, C++, Java