

## SELECT WORK EXPERIENCE

---

### Hinge

#### Staff Machine Learning Engineer

New York, NY, Apr 2024 – Current

*Leading the democratization and platformization of AI at Hinge.*

- Spearheading a workstream on **online monitoring and evaluation of generative AI systems**. Leading vendor evaluations for solutions to monitoring generative AI systems, prompt management, A/B testing, etc. Coordinating initial platformization of disparate generative AI initiatives across backend, AI, and data engineering organizations to facilitate structured evaluation and measurement. Defining the strategy and roadmap for judge-based evaluations, tracing of production foundation model invocations via OpenTelemetry, and hot-path guardrail of LLM inputs and outputs.
- Owned the **model serving strategy** from end-to-end, delivering a generally-available platform in three quarters that reduced time-to-production for net-new models by ~50%. Developed year-long initial strategy from first principles, aligning the AI Platform Core team on core initial focuses – system and model observability; model governance; unified artifact management for heterogeneous deployment environments. Drove close collaboration across cross-functional platform partner teams in Data and Engineering.
- Established the **Hinge AI Staff Engineer Forum**, providing a network of six Staff-level technical leaders across AI at Hinge with a venue for advising AI leadership on departmental technical strategy, as well as for sharing tips regarding “leadership without authority,” particularly with regards to delegation, mentorship and coaching.

### Spotify

#### Senior Machine Learning Engineer

New York, NY, Mar 2021 – Mar 2024

*Senior member of the ML Platform org. Primarily focused on managed multi-tenant compute infrastructure for KubeFlow and Ray in GCP. Later shifted focus to AI/ML governance and artifact/experiment management solutions.*

- Drove early-stage evaluation and implementation of centralized, multi-tenant Ray infrastructure based on KubeRay;
- Spearheaded development of Backstage-based model governance solution in preparation for the EU Digital Services Act;
- Defined and drove general-availability roadmap for KubeFlow-based orchestration product; defined SLOs and educated ML Platform org on SRE best practices.

### NVIDIA

#### Senior Systems Software Engineer, AI Infrastructure

New York, NY, Dec 2019 – Jan 2021

*Worked on AI infrastructure for autonomous vehicle research and development.*

### Twitter

#### Machine Learning Software Engineer

New York, NY, Aug 2018 – Dec 2019

*Member of Cortex, Twitter’s central ML platform organization. Worked on: workflow orchestration; experiment management.*

## EDUCATION

---

University of Chicago B.S. Computer Science, B.A Economics

2011 – 2015

## SPEAKING

---

Data Council, “Trimming the Long Tail of Production Model Ownership at Hinge”

Apr 2025

Ray Summit, “Accelerating Time-to-Production for ML at Hinge”

Sep 2024

MLconf, “Empowering Traceable and Auditable ML in Production at Spotify with Hendrix”

Mar 2022

KubeCon + CloudNativeCon, “Scaling KubeFlow for Multi-tenancy at Spotify”

Oct 2021

## SKILLS

---

- **Machine Learning and Data:** Ray, PyTorch, Apache Spark, Apache Flink, MLflow, Apache Airflow, KubeFlow, TensorFlow
- **Infrastructure and Distributed Systems:** Kubernetes, Terraform, Helm, Prometheus, OpenTelemetry, gRPC, Docker
- **Enterprise Platforms:** Amazon Web Services (AWS), Google Cloud (GCP), Databricks
- **Programming Languages:** Python, Go, C++, Java