

SELECT WORK EXPERIENCE

Hinge

Staff Machine Learning Engineer

New York, NY, Apr 2024 – Current

Leading the democratization and platformization of AI at Hinge.

- Spearheading **online monitoring and evaluation of generative AI systems**. Leading vendor evaluations for solutions to monitoring generative AI systems, prompt management, A/B testing, etc. Implemented a platform for generative AI systems, prompt management, and judge-based evaluations based on OpenTelemetry, [OpenInference](#), and [Phoenix](#) in collaboration with stakeholders across Backend, AI, and Data disciplines. Defining and steering the strategy, roadmap, and platform architecture for judge-based evaluations, tracing of production foundation model invocations via OpenTelemetry, and hot-path guardrail of LLM inputs and outputs.
- Owned the **model serving strategy** from end-to-end. Delivered a generally-available platform based on MLflow, Ray, and Spark in three quarters that reduced time-to-production for net-new models by ~50%. Developed year-long strategy and roadmap from first principles, aligning the AI Platform Core team on core focuses – system and model observability; model governance; unified artifact management for heterogeneous deployment environments. Drove close collaboration across cross-functional platform partner teams in Data and Engineering.
- Established the **Hinge AI Staff Engineer Forum**, providing a network of six Staff-level technical leaders across AI at Hinge with a venue for advising AI leadership on departmental technical strategy, as well as for sharing tips regarding “leadership without authority,” particularly with regards to delegation, mentorship and coaching.

Spotify

Senior Machine Learning Engineer

New York, NY, Mar 2021 – Mar 2024

Senior member of the ML Platform org. Primarily focused on managed multi-tenant compute infrastructure for Kubeflow and Ray in GCP. Later shifted focus to AI/ML governance and artifact/experiment management solutions.

- Drove early-stage evaluation and implementation of centralized, multi-tenant Ray infrastructure based on KubeRay;
- Spearheaded development of Backstage-based model governance solution in preparation for the EU Digital Services Act;
- Defined and drove general-availability roadmap for Kubeflow-based orchestration product; defined SLOs and educated ML Platform org on SRE best practices.

NVIDIA

Senior Systems Software Engineer, AI Infrastructure

New York, NY, Dec 2019 – Jan 2021

Worked on AI infrastructure for autonomous vehicle research and development.

Twitter

Machine Learning Software Engineer

New York, NY, Aug 2018 – Dec 2019

Member of Cortex, Twitter’s central ML platform organization. Worked on: workflow orchestration; experiment management.

EDUCATION

University of Chicago B.S. Computer Science, B.A Economics

2011 – 2015

SPEAKING

Data Council, “[Trimming the Long Tail of Production Model Ownership at Hinge](#)”

Apr 2025

Ray Summit, “[Accelerating Time-to-Production for ML at Hinge](#)”

Sep 2024

MLconf, “[Empowering Traceable and Auditable ML in Production at Spotify with Hendrix](#)”

Mar 2022

KubeCon + CloudNativeCon, “[Scaling Kubeflow for Multi-tenancy at Spotify](#)”

Oct 2021

SKILLS

- **Machine Learning and Data:** [Ray](#), [PyTorch](#), [Apache Spark](#), [Apache Flink](#), [MLflow](#), [Apache Airflow](#), [Kubeflow](#), [TensorFlow](#)
- **Infrastructure and Distributed Systems:** Kubernetes, Terraform, Helm, Prometheus, OpenTelemetry, gRPC, Docker
- **Enterprise Platforms:** Amazon Web Services (AWS), Google Cloud (GCP), [Databricks](#)
- **Programming Languages:** Python, Go, C++, Java