

## SELECT WORK EXPERIENCE

---

### Hinge

#### Staff Machine Learning Engineer, AI Platform

New York, NY, Apr 2024 – Current

*Leading the platformization of Hinge's various AI/ML efforts (recommendations, trust and safety, user growth and retention, etc.).*

- Defined the org-wide technical direction and strategy for production serving of ML models at Hinge. Conducted in-depth and methodical evaluation of two solutions (Ray Serve and [Databricks Model Serving](#)). Developed multi-quarter development roadmap centered around top priorities (system and model observability; model governance; structured rollout and deployments; etc.); drove subsequent alignment and collaboration between AI Platform team and cross-functional partners, e.g. Data Platform and Engineering Platform teams.
- Re-architected photo-scoring pipeline for [Top Photo](#) to enable multi-plexed rescoring of user photos in near-real-time using Databricks Delta Live Tables, Spark, Flink, and Ray Serve; reduced rollout time of new photo-scoring models from months to one week.

### Spotify

#### Senior Machine Learning Engineer

New York, NY, Mar 2021 – Mar 2024

*Senior member of the ML Platform org. Primarily focused on managed multi-tenant compute infrastructure for Kubeflow and Ray in GCP. Later shifted focus to AI/ML governance and artifact/experiment management solutions.*

- Drove early-stage evaluation and implementation of centralized, multi-tenant Ray infrastructure based on KubeRay;
- Spearheaded development of Backstage-based model governance solution in preparation for the EU Digital Services Act;
- Defined and drove general-availability roadmap for Kubeflow-based orchestration product; defined SLOs and educated ML Platform org on SRE best practices.

### NVIDIA

#### Senior Systems Software Engineer, AI Infrastructure

New York, NY, Dec 2019 – Jan 2021

*Worked on AI infrastructure for autonomous vehicle research and development.*

### Twitter

#### Machine Learning Software Engineer

New York, NY, Aug 2018 – Dec 2019

*Member of Cortex, Twitter's central ML platform organization. Worked on: workflow orchestration; experiment management.*

## EDUCATION

---

University of Chicago B.S. Computer Science, B.A Economics

2011 – 2015

## SPEAKING

---

Data Council, "[Trimming the Long Tail of Production Model Ownership at Hinge](#)"

Apr 2025

Ray Summit, "[Accelerating Time-to-Production for ML at Hinge](#)"

Sep 2024

MLconf, "[Empowering Traceable and Auditable ML in Production at Spotify with Hendrix](#)"

Mar 2022

KubeCon + CloudNativeCon, "[Scaling Kubeflow for Multi-tenancy at Spotify](#)"

Oct 2021

## SKILLS

---

- **Machine Learning and Data:** [Ray](#), [PyTorch](#), [Apache Spark](#), [Apache Flink](#), [MLflow](#), [Apache Airflow](#), [Kubeflow](#), [TensorFlow](#)
- **Infrastructure and Distributed Systems:** Kubernetes, Terraform, Prometheus, gRPC, Docker
- **Enterprise Platforms:** Amazon Web Services (AWS), Google Cloud (GCP), [Databricks](#)
- **Programming Languages:** Python, Go, C++, Java